

# The Normal Distribution

*The second statistical distribution we are going to look at is the normal or Gaussian distribution. This distribution is of central importance to nearly all areas of statistics, and especially those areas dealing with statistical analysis. This means a good understanding of the normal distribution useful in many different areas.*

## Contents

|                                                                    |    |
|--------------------------------------------------------------------|----|
| Introduction . . . . .                                             | 2  |
| Aim . . . . .                                                      | 2  |
| Learning Outcomes . . . . .                                        | 2  |
| Defining the Normal Distribution . . . . .                         | 2  |
| General Definition . . . . .                                       | 2  |
| The Standard Normal Distribution . . . . .                         | 3  |
| Quantiles and Distribution Estimates . . . . .                     | 4  |
| Normal Distribution in SciPy . . . . .                             | 4  |
| The Standard Normal Curves . . . . .                               | 4  |
| Creating an Interactive PDF for the Normal Distribution . . . . .  | 7  |
| Extra Mile: Non-Standard Normal Distributions . . . . .            | 12 |
| Quantiles and Probabilities . . . . .                              | 12 |
| Calculating Probabilities from the Standard Normal Curve . . . . . | 12 |
| Using Sub-Figures to Plot Probabilities . . . . .                  | 13 |
| Questions . . . . .                                                | 18 |
| References . . . . .                                               | 19 |

## Introduction

### Aim

The second statistical distribution we are going to look at is the *normal* or *Gaussian* distribution. This distribution is of central importance to nearly all areas of statistics, and especially those areas dealing with statistical analysis. This means a good understanding of the normal distribution useful in many different areas.

Part of the reason for the significance of normal distribution lies in the assumption that *many* 'naturally occurring populations follow this distribution (Urdan, 2022b). For instance, the demographics of most human or animal populations, errors and quality estimates, etc. This makes the normal distribution the first distribution that we often turn to when simulating natural processes.

The second reason for the importance of the normal distribution is that *most* common statistical descriptors rely on the normal distribution as a reference<sup>1</sup> (Urdan, 2022a). For instance the distribution of statistical errors is usually assumed to be normal, and many of the population and sample descriptors we use are also based on the normal distribution. We have looked at some of those descriptors already in *Workshop W.00 (The Uniform Distribution)*, such as the *mean* and *variance*. By studying the normal distribution, we will see where many of these concepts and terms come from.

### Learning Outcomes

- To understand the common measures of central tendency, and how these relate to the normal distribution
- To understand the common measures of variability (variance), and how these relate to the normal distribution
- To understand the definition of the normal distribution, and how this definition is used
- To use NumPy and SciPy to calculate and apply basic properties of the normal distribution

## Defining the Normal Distribution

### General Definition

The second statistical distribution we are going to look at is the Normal or Gaussian distribution which only has a continuous form.

---

<sup>1</sup>See *Explanation EC.01 (From Data to Distribution)*, *EC.02 (Samples and Populations)* and *EC.03 (Hypothesis Testing)* for more details of how the normal distribution is used in statistical analysis.

**Table 1:** Definitions of the Normal Distribution

| Term                               | Expression                                                                                                                                                                                                                                                                        |
|------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Mean                               | $\mu$                                                                                                                                                                                                                                                                             |
| Variance                           | $\sigma^2$                                                                                                                                                                                                                                                                        |
| Probability Density Function (PDF) | $f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$                                                                                                                                                                                                           |
| Cumulative Density Function (CDF)  | $F(x) = \Phi\left(\frac{x-\mu}{\sigma}\right)$<br>$= \frac{1}{2} \left[ 1 + \operatorname{erf}\left(\frac{x-\mu}{\sigma\sqrt{2}}\right) \right]$<br>where $\operatorname{erf}(x) = \frac{1}{\sqrt{\pi}} \int_{-x}^x e^{-t^2} dt$<br>$= \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$ |

Note that the *mean*,  $\mu$ , and *variance*,  $\sigma^2$ , are parameters of the normal distribution. This will become important later when we talk about the descriptive statistics used for characterising populations<sup>2</sup>. Also remember from *Workshop W.00 (The Uniform Distribution)* that the *standard deviation*, *sigma* is defined as

$$\sigma = \sqrt{\sigma^2}$$

This is important as we often use the standard deviation, and not the variance directly, to define specific normal distributions. Knowing which measure you are using is important to make sure you are drawing the right curve...

The *error function* erf in Table 1 is also a concept you will see in most definitions of methods used for statistical hypothesis testing<sup>3</sup>. Many of the estimates of *statistical confidence* are based around the normal definition, and the erf function will turn up regularly in these discussions.

### The Standard Normal Distribution

An important variation of the normal distribution is the *standard normal distribution*, which is used to define numerous statistical concepts. The standard normal distribution occurs when  $\mu = 0$  and  $\sigma^2 = 1$ , giving a PDF of

$$\varphi(z) = \frac{e^{-z^2/2}}{\sqrt{2\pi}}$$

<sup>2</sup>See *Explanation EC.01 (From Data to Distribution)* and *EC.02 (Samples and Populations)* for details.

<sup>3</sup>See *Explanation EC.03 (Hypothesis Testing)* for some of the tools and methods this applies to

and a CDF of

$$\begin{aligned}\Phi(z) &= \frac{1}{\sqrt{2\pi}} \int_{-\infty}^z e^{-t^2/2} dt \\ &= \frac{1}{2} \left[ 1 + \operatorname{erf}\left(\frac{z}{\sqrt{2}}\right) \right]\end{aligned}$$

When you see ‘ $\varphi$ ’ or ‘ $\Phi$ ’ in statistical tests or definitions (such as Table 1) think of the standard normal distributions. You will also see  $\Phi^{-1}$ , defined for  $z$  in the interval  $z \in (0, 1)$  as

$$\Phi^{-1}(z) = \sqrt{2} \operatorname{erf}^{-1}(z_p - 1)$$

### Quantiles and Distribution Estimates

The inverse function  $\Phi^{-1}$  is often used in *quantile analysis*. A *quantile* is like a contour line dividing a probability range, or grouping related probabilities. We use the quantile of a normal distribution to determine the *confidence* we have that a specific value, typically the mean, falls inside or outside a probability range.

For instance the quantile  $z_p$  is often used as a confidence estimate for a specific probability  $p$  by (Wichura, 1988)

$$z_p = \Phi^{-1}(p)$$

A random variable  $X$  taken from a normal distribution with a mean of  $\mu$  will exceed  $\mu + z_p\sigma$  with probability  $1 - p$ , and will lie outside the interval  $\mu \pm z_p\sigma$  with probability  $2(1 - p)$  (Statistics LibreTexts, 2020). Since  $z_p$  is continuous, estimates of  $\Phi^{-1}$  are commonly used (Hyndman and Fan, 1996) in most cases<sup>4</sup>. In SciPy estimates of  $z_p$  can use gained from the *percentage point function (ppf)*.

For any normal distribution around 68% of the values are within one standard deviation  $\sigma$  of from the mean; about 95 % of values are within two standard deviations; and about 99.7% of the values are within 3 standard deviations of the mean.

## Normal Distribution in SciPy

### The Standard Normal Curves

We have covered the details of calculating the PDF and CDF of a distribution in *Workshop W.00 (The Uniform Distribution)*. Since SciPy has a standardised interface for most common distributions, the process for the normal distribution is much the same as we used previously.

Here the standard class is `scipy.stats.norm`, and we start with specifying `loc` and `scale` as before. For the `scipy.stats.norm` class, `loc` specifies the mean of the normal curve,

---

<sup>4</sup>See Reference RM.01 (*Inferential Statistics*) for a table of common estimates of  $z_p$ . How these estimates are used to determine statistical confidence is covered in Explanation EC.02 (*Samples and Populations*).

and `scale` is the standard deviation. We will also define a mean and `stdev`, and since this is a standard normal curve  $\mu = 0$  and  $\sigma = 1$ . This gives

```
1 mean = 0
2 stdev = 1
3
4 loc = mean
5 scale = stdev
```

Our normal distribution can then be generated by<sup>5</sup>

```
1 normal_distribution = sp.stats.norm(loc, scale)
```

We can ask for the mean and variance of `normal_distribution` using the standard method call to `stats()`, as we did for the uniform distribution

```
1 stat_mean, stat_var = normal_distribution.stats(
2     moments="mv"
3 )
4
5 print(
6     "Standard normal distribution\n"
7     f"Mean: {stat_mean:.2f}\nVariance: {stat_var:.4f}\n"
8 )
```

```
Standard normal distribution
Mean: 0.00
Variance: 1.0000
```


We will display the standard normal curve as  $\pm 5\sigma$  of the mean. So we will need the NumPy `linspace()` method to give the appropriate range for the  $x$  axis

```
1 # Generate the values for the x-axis between -5.0 and 5.0
2 # standard deviations
3 normal_x_values = np.linspace(-5.0, 5.0, 100000)
```

Now we can display the PDF of the `normal_distribution` we have just created in Figure 1

```
1 # The Plot of the Standard Normal PDF where  $\mu = 10$ 
2 # and  $\sigma = 1$  in the Real Interval  $[-5.0, 5.0]$ 
3 # fig:W01_StdNormPDF1
4
5 x_axis = normal_x_values
6 y_axis = normal_distribution.pdf(normal_x_values)
7
```

<sup>5</sup>This assumes the *standard header* used in the Workbooks has already been loaded. See *HowTo H.00 (New Jupyter Workbooks)* for the code you will need in a blank Jupyter workbook.



**Figure 1:** The Plot of the Standard Normal PDF where  $\mu = 10$  and  $\sigma = 1$  in the Real Interval  $[-5.0, 5.0]$

```

8      # Plot the result
9      plot.figure(figsize=(8, 4))
10
11     # Format the plot
12     plot.grid(True, which="both")
13
14     # Label the axes
15     plot.xlabel("$\\sigma$")
16     plot.ylabel("f(x)")
17
18     # Label the plot
19     plot.title(
20         "Probability Density Function of the Standard Distribution\n"
21         "Where  $\mu =$ "
22         + str(mean)
23         + " and  $\sigma =$ "
24         + str(stdev),
25         pad=24.0,
26     )
27
28     # Display the graph
29     plot.plot(x_axis, y_axis)


```

As before, displaying the CDF is essentially identical to the PDF, apart from the change in Line 6. This should give the standard CDF plot in Figure 2.

```

1      # The Plot of the Standard Normal CDF where  $\mu = 10$ 
2      # and  $\sigma = 1$  in the Real Interval  $[-5.0, 5.0]$ 
3      # fig:W01_StdNormCDF1
4

```



**Figure 2:** The Plot of the Standard Normal CDF where  $\mu = 0$  and  $\sigma = 1$  in the Real Interval  $[-5.0, 5.0]$

```

5  x_axis = normal_x_values
6  y_axis = normal_distribution.cdf(normal_x_values)
7
8  # Plot the result
9  plot.figure(figsize=(8, 4))
10
11 # Format the plot
12 plot.grid(True, which="both")
13
14 # Label the axes
15 plot.xlabel("$\\sigma$")
16 plot.ylabel("F(x)")
17
18 # Label the plot
19 plot.title(
20     "Cumulative Density Function of the Standard Distribution\n"
21     "Where $\\mu = $"
22     + str(mean)
23     + " and $\\sigma = $"
24     + str(stdev),
25     pad=24.0,
26 )
27
28 # Display the graph
29 plot.plot(x_axis, y_axis)

```

### Creating an Interactive PDF for the Normal Distribution

To explore other values of the normal distribution, outside the standard normal distribution, we will go beyond a static plot. In Jupyter you can animate different variables, using the

**Jupyter Widgets** library. This requires Javascript support, which means it is usually confined to a browser: but Visual Studio Code can cope with these as well.

Import the widgets library as

```
1 | # Import the widget packages for this Notebook
2 | import ipywidgets as widgets
```

Now the code for the interactive normal distribution can be created as follows. This is very similar to the charts we have created before, except that the call to plot is wrapped in a function called `plot_normal_pdf` (Line 6). The `plot_normal_pdf` takes the parameters `mean` and `stdev`, and calculates the PDF using these values. To show the spread of the normal distribution, `plot_normal_pdf` also annotates the mean and first standard deviation in Line 37 to Line 58. Then Line 64 creates a variable `run_plot` from the Jupyter `widgets.interact()` method, and asking for two sliders to represent the mean and `stdev`. Finally Line 79 calls `run_plot`, which should show the slides and allow you to modify the plot values.

```
1 | ##
2 | ## Define the Interactive Function for the Normal PDF
3 | ##
4 |
5 |
6 | def plot_normal_pdf(mean, stdev):
7 |     # Plots the Normal Distribution of with a mean of `mean`,
8 |     # and a standard deviation of `stdev`
9 |
10 |     # Calculate the interval between 0 and 20, in steps of 0.001
11 |     x_axis = np.arange(0, 20, 0.001)
12 |     y_axis = sp.stats.norm.pdf(x_axis, mean, stdev)
13 |
14 |     # Plot the result
15 |     plot.figure(figsize=(8, 4))
16 |     plot.plot(x_axis, y_axis)
17 |
18 |     plot.grid(True, which="both")
19 |
20 |     plot.title(
21 |         "Probability Density Function of the Standard"
22 |         " Normal Distribution With\n a Mean of "
23 |         + str(mean)
24 |         + " and Standard Deviation of "
25 |         + str(stdev),
26 |         pad=24.0,
27 |     )
28 |     plot.xlabel("Value")
29 |     plot.ylabel("Relative Likelihood (p)")
30 |
31 |     plot.ylim(
32 |         0,
33 |         sp.stats.norm.pdf(mean, mean, stdev) * 1.08,
```



```

34     )
35     plot.xlim(0, 20)
36
37     # Annotate Mean and Standard Deviations
38     plot.vlines(
39         mean,
40         0,
41         sp.stats.norm.pdf(mean, mean, stdev),
42         color=signal_colours["orange"],
43         linestyle="dashed",
44     )
45     plot.vlines(
46         mean - stdev,
47         0,
48         sp.stats.norm.pdf(mean - stdev, mean, stdev),
49         color=signal_colours["orange"],
50         linestyle="dotted",
51     )
52     plot.vlines(
53         mean + stdev,
54         0,
55         sp.stats.norm.pdf(mean + stdev, mean, stdev),
56         color=signal_colours["orange"],
57         linestyle="dotted",
58     )
59
60     plot.show()
61
62     # Run the interactive plot
63     run_plot = widgets.interact(
64         plot_normal_pdf,
65         mean=widgets.FloatSlider(
66             description="Mean",
67             min=0,
68             max=20,
69             value=10,
70         ),
71         stdev=widgets.FloatSlider(
72             description="Standard Deviation",
73             min=1,
74             max=5,
75             value=2,
76         ),
77     ),
78 )
79
80 run_plot

```

```

1     {
2         "model_id": "0c312dad5c9241eb8b77c93bc86e9e94",

```

```

3     "version_major": 2,
4     "version_minor": 0
5 }

```

The interactive version of the CDF plot is essentially identical, apart from the changing the SciPy library calls. Notice that for large standard deviations, or for values of the mean close to the limit, the CDF curve doesn't fit into the 'window' of values we have defined. You can also see how changing the standard deviation causes the plot to 'rotate' about the mean. Whilst changing the mean changes the location of the curve, just as for the PDF plot above.

```

1  ##
2  ## Define the Interactive Function for the Normal CDF
3  ##
4
5
6  def plot_normal_cdf(mean, stdev):
7      # Plots the Normal Distribution of with a mean of `mean`,
8      # and a standard deviation of `stdev`
9
10     # Calculate the interval between 0 and 20, in steps of 0.001
11     x_axis = np.arange(0, 20, 0.001)
12     y_axis = sp.stats.norm.cdf(x_axis, mean, stdev)
13
14     # Plot the result
15     plot.figure(figsize=(8, 4))
16     plot.plot(x_axis, y_axis)
17
18     plot.grid(True, which="both")
19
20     plot.title(
21         "Cumulative Density Function of the Standard"
22         " Normal Distribution With\n a Mean of "
23         + str(mean)
24         + " and Standard Deviation of "
25         + str(stdev),
26         pad=24.0,
27     )
28     plot.xlabel("Value")
29     plot.ylabel("Relative Likelihood (p)")
30
31     plot.ylim(0, 1.1)
32     plot.xlim(0, 20)
33
34     # Annotate Mean and Standard Deviations
35     plot.vlines(
36         mean,
37         0,
38         sp.stats.norm.cdf(mean, mean, stdev),
39         color=signal_colours["orange"],
40         linestyle="dashed",

```

```

41     )
42     plot.vlines(
43         mean - stdev,
44         0,
45         sp.stats.norm.cdf(mean - stdev, mean, stdev),
46         color=signal_colours["orange"],
47         linestyle="dotted",
48     )
49     plot.vlines(
50         mean + stdev,
51         0,
52         sp.stats.norm.cdf(mean + stdev, mean, stdev),
53         color=signal_colours["orange"],
54         linestyle="dotted",
55     )
56
57     plot.show()
58
59     # Run the interactive plot
60     run_plot = widgets.interact(
61         plot_normal_cdf,
62         mean=widgets.FloatSlider(
63             description="Mean",
64             min=0,
65             max=20,
66             value=10,
67         ),
68         stdev=widgets.FloatSlider(
69             description="Standard Deviation",
70             min=1,
71             max=5,
72             value=2,
73         ),
74     ),
75 )
76
77 run_plot

```

```

1  {
2      "model_id": "7aa84250203b47da910795ef1d9a95d8",
3      "version_major": 2,
4      "version_minor": 0
5  }

```

You can use the sliders to explore both the simple normal distribution, and observe the effect changing the mean and standard deviation has. Note that in our model, the central value *always* follows the mean (dashed line), and the standard deviation is *always* symmetric (dotted lines). You can check this by observing the shapes of the PDF and CDF as it changes.

### Extra Mile: Non-Standard Normal Distributions

In ‘real’ normal distribution we sometimes find the ‘central’ value is *not* the mean, but is shifted from the mean. This is called *skew*, and the resulting distributions are *non-central*. You can explore the effect skew has using the `scipy.stats.skewnorm` distribution instead of `scipy.stats.norm` in the code above. Some normal distributions are also ‘fatter’ or ‘thinner’ than you would expect, and this is called *kurtosis*. For a given set of data, SciPy can calculate kurtosis using the `scipy.stats.kurtosis()` method. Similarly the degree of skew can be calculated using `scipy.stats.skew()`.

Try modifying the code above to represent a wider range of normal distributions for both the PDF and CDF. You can start with the `scipy.stats.skewnorm` distribution, and add a slider to represent the skew. You will probably also have to modify the limits used by Matplotlib, either dynamically or at least with a larger window. The ‘cheatsheets’ for Matplotlib are a good place to start when modifying plots.

### Quantiles and Probabilities

#### Calculating Probabilities from the Standard Normal Curve

As we said earlier, a core use of the normal distribution in statistical analysis is to calculate *quantiles*. These ‘probability regions’ of the normal distribution are used to model both real-world populations, and serve as the core theory for a number of statistical methods. We will look at both uses later<sup>6</sup>, for the moment focusing on the calculations themselves.

In general, though, we are interested in two different types of regions for a normal distribution

- **one-tailed** probabilities address questions where a quantity is either *above* or *below* a specified limit on the left **or** the right of the normal distribution
- **two-tailed** probabilities address questions where a quantity is either *between* two limits, or *outside* those limits on **both** sides of the normal distribution

Knowing the question you want to answer is important for selecting the correct region. For example, some statistical tests are one-tailed, and so *cannot* address problems where we need to find whether a specific probability lies in a given range on the ‘other side’ of the distribution. Other tests come in both one-tailed and two-tailed forms, and again, selecting the right kind of statistical test is important for getting an answer that has some validity.

For the following examples, we will use  $x$  as the variable of interest in the probability range  $P$ , where  $P \in (0, 1]$  (since the normal curve is asymptotic at  $P = 0$ ). Specifically we are going to look at two-tailed regions where

1.  $P(x \leq 0.4)$
2.  $P(x \geq -1.3)$
3.  $P(-0.1 \leq x \leq 0.4)$
4.  $P(|x| \leq 1.2) = P(-1.2 \leq x \leq 1.2)$

---

<sup>6</sup>See *Explanation EC.03 (Hypothesis Testing)* for more uses of the normal distribution in both real-world and statistical modelling. *Explanation EC.02 (Samples and Populations)* covers the use of the normal distribution for statistical qualification, especially in the analysis of error and statistical confidence.

Calculating  $P$  for these probabilities can be done using the `scipy.stats.norm.cdf()` we have just used in the section above. Before we have used the `cdf()` method by passing a range of  $x$  values: either `normal_x_values` or the `np.arange()` method in the interactive plots. However we can also pass in a single value, and SciPy will return the probability calculation for the specified point on the CDF curve. If we look at the definition of the `cdf` method we will also find the parameters are defined as

```
1 | cdf(x, loc=0, scale=1)]
```

Since we are looking for probabilities on the *standard normal* curve, we don't need to specify either `loc` or `scale`. If we needed to, though, we could adjust the probability calculation for the specific mean and standard deviation of the population in questions.

For our calculations, though, the four probabilities we need are given by

```
1 | # $P(x \leq 0.4)$
2 | p1 = sp.stats.norm.cdf(0.4)
3 | # $P(x \geq -1.3)$
4 | p2 = 1 - sp.stats.norm.cdf(-1.3)
5 | # $P(-0.1 \leq x \leq 0.4)$
6 | p3 = p1 - sp.stats.norm.cdf(-0.1)
7 | # $P(|x| \leq 1.2) = P(-1.2 \leq x \leq 1.2)$
8 | p4 = sp.stats.norm.cdf(1.2) - sp.stats.norm.cdf(-1.2)
9 |
10 | print(f"p1: {p1 :.3f}")
11 | print(f"p2: {p2 :.3f}")
12 | print(f"p3: {p3 :.3f}")
13 | print(f"p4: {p4 :.3f}")
```

```
p1: 0.655
p2: 0.903
p3: 0.195
p4: 0.770
```

### Using Sub-Figures to Plot Probabilities

Plotting these probabilities requires first determining the region of interest. Before we looked at the interval  $\pm 5\sigma$ , but we will restrict this slightly to  $\pm 3\sigma$  for the four regions


```
1 | # Generate the values for the x-axis between -3.0 and 3.0
2 | # standard deviations
3 | normal_x_values = np.linspace(-3.0, 3.0, 100000)
```

Next we need the PDF of the standard normal distribution, calculated over the  $x$  axis range, `normal_x_values`, as before. Since this is the standard normal distribution, we can ignore `loc` and `scale`

```
1 | # Create probability density function (PDF)
2 | std_normal_pdf = sp.stats.norm.pdf(normal_x_values)
```

We want to plot the four plots together, so we will use a feature of Matplotlib called **subplots**. Here figure will hold the final plot image, and subfig will hold the individual plots. The individual plots will behave in exactly the same way as before: all we are doing here is effectively creating an array of them

```
1 | figure, subfig = plot.subplots(
2 |     ncols=2,
3 |     nrows=2,
4 |     layout="constrained",
5 | )
```



To the blank set of plots, we now want to add the standard normal curves. This uses the `normal_x_values` and `std_normal_pdf` values we have just calculated. Since we are walking over the 2x2 array, we will also use this opportunity to set x and y axis labels for the left-most and bottom plots. We are using the `constrained` layout to save space in the figure, so this avoids having to add axis values where we don't need them (and where it is still clear what the plot represents)

```
1 | for row in range(2):
2 |     for col in range(2):
3 |         subfig[row, col].plot(
4 |             normal_x_values, std_normal_pdf
5 |         )
6 |         if row == 1:
7 |             subfig[row, col].set_xlabel("Values")
```

```

8         if col == 0:
9             subfig[row, col].set_ylabel(
10                 "Probability Density"
11             )

```

With the standard normal distributions in place, we need to calculate (and plot) the regions of interest. To do this we will iterate back through the 2x2 array of figures, and set-up each sub-figure again.

We have already seen the `set_title()` method for other plots, as the `title` method. Each sub-plot behaves as an independent axis in figure, and is configured via the `Axes class`<sup>7</sup>. So we have to use the methods of the `Axes` class to add a title to show the region of interest (Line 6, Line 17, Line 28 and Line 40). This is why we also have to use `matplotlib.axes.Axes.set_ylim()` method in Line 14, instead of the normal `ylim()` method we used previously. Both do the same job, but `set_ylim()` applies to sub-figures and not figures.

Of more interest is the `matplotlib.pyplot.fill_between()` method. The first arguments of `fill_between()` are the values of the  $x$  and  $y$  axes; which by default will just colour in the full plot. However we can also pass in an array of Boolean arguments as in Line 11, Line 22, Line 33--34 and Line 45--46. These arguments tell Matplotlib to *only* colour the chart if the specified conditions are true. So we can now add the two probabilities, and two probability ranges, and we should get coloured regions shown in Figure 3.

```


1  # A Plot of Four Quantiles, Showing the Regions Defined by
2  # Four Probabilities on the Standard Normal PDF
3  # fig:W01_StdNormQuantile
4
5  # Create the sub-figure plot for p1
6  subfig[0, 0].set_title(f"P(x ≤ 0.4) = {p1: .3f}")
7
8  subfig[0, 0].fill_between(
9      normal_x_values,
10     std_normal_pdf,
11     alpha=0.2,
12     where=(normal_x_values ≤ 0.4),
13 )
14
15 subfig[0, 0].set_ylim(0, 0.5)
16
17 # Create the sub-figure plot for p2
18 subfig[0, 1].set_title(f"P(x ≥ -1.3) = {p2: .3f}")
19
20 subfig[0, 1].fill_between(
21     normal_x_values,
22     std_normal_pdf,
23     alpha=0.2,

```

<sup>7</sup>The class name chosen by Matplotlib for sub-figures isn't that intuitive, and so we have chosen 'subfig' as the object name referring to 'regions of figure'. In many examples, you will see these 'sub-figures' given object names such as 'axes', or 'axs' after the class name. They all refer to the same `Axes` class.

```
24     where=(normal_x_values >= -1.3),
25 )
26
27 subfig[0, 1].set_ylim(0, 0.5)
28
29 # Create the sub-figure plot for p3
30 subfig[1, 0].set_title(f"P(-0.1 ≤ x ≤ 0.4) = {p3: .3f}")
31
32 subfig[1, 0].fill_between(
33     normal_x_values,
34     std_normal_pdf,
35     alpha=0.2,
36     where=(normal_x_values <= 0.4)
37     & (normal_x_values >= -0.1),
38 )
39
40 subfig[1, 0].set_ylim(0, 0.5)
41
42 # Create the sub-figure plot for p4
43 subfig[1, 1].set_title(f"P(|x| ≤ 1.2) = {p4: .3f}")
44
45 subfig[1, 1].fill_between(
46     normal_x_values,
47     std_normal_pdf,
48     alpha=0.2,
49     where=(normal_x_values <= 1.2)
50     & (normal_x_values >= -1.2),
51 )
52
53 subfig[1, 1].set_ylim(0, 0.5)
54
55 # Plot the final figure
56 figure
```





**Figure 3:** A Plot of Four Quantiles, Showing the Regions Defined by Four Probabilities on the Standard Normal PDF

## Questions

Use Python to answer the following questions. Where appropriate give answers to 4 decimal places.

1. Given a normal distribution with  $\mu = 10$  and  $\sigma = 2$ 
  - a. Plot the quantiles and  $P$  values for the following regions
    - $P(|x| \leq \sigma)$
    - $P(|x| \leq 2\sigma)$
    - $P(|x| \leq 3\sigma)$
    - $P(\mu - \sigma \leq x \leq \mu + \sigma)$
  - b. How do the  $P$  values above relate to the expected values for the area under a standard normal curve?

Answer to 1(a):

```
1 | # code for question 1(a)
```

Answer to 1(b)

<!-- Markdown here -->

2. The lifetime  $T$  of a particular model of battery is normally distributed with an average lifetime of 825 hours and a variance of 900 hours. Using Python can you answer the following - illustrating your answers by shading the pdf:
  - a. What is the probability that a battery lasts at least 900 hrs?
  - b. What is the probability that a battery last less than 800 hours?
  - c. What is the probability that the lifetime of a battery is between 810 and 870 hours?

```
1 | # code for question 2(a) -- 2(c)
```

3. The resistance of a 10% 1k $\Omega$  is modelled using a normal distribution where the 10% is assumed to 95% of the distribution.
  - a. What is the standard deviation for the model?
  - b. Using Python, what is the probability that the measured resistance  $R$  will be less than 950 $\Omega$ ?
  - c. What is the probability  $R \geq 1.1\text{k}\Omega$ ?
  - d. What is the probability  $R$  will be between 0.98 & 1.02k $\Omega$ ?

Answer:

```
1 | # code for question 3(b) -- 3(d)
```

---

## References

3.6: Distribution and Quantile Functions (2020) [Online]. Available from: <[https://stats.libretexts.org/Bookshelves/Probability\\_Theory/Probability\\_Mathematical\\_Statistics\\_and\\_Stochastic\\_Processes\\_\(Siegrist\)/03%3A\\_Distributions/3.06%3A\\_Distribution\\_and\\_Quantile\\_Functions](https://stats.libretexts.org/Bookshelves/Probability_Theory/Probability_Mathematical_Statistics_and_Stochastic_Processes_(Siegrist)/03%3A_Distributions/3.06%3A_Distribution_and_Quantile_Functions)> [Accessed 4 June 2025].

Hyndman, R. J. and Fan, Y. (1996) Sample Quantiles in Statistical Packages. *The American Statistician* [Online], 50 (4), pp. 361–365. Available from: <<https://www.jstor.org/stable/2684934>> [Accessed 4 June 2025].

Urdan, T. C. (2022a) Statistical Significance, Effect Size, and Confidence Intervals [Online]. In: *Statistics in Plain English*. Routledge, pp. Chapter 7. Available from: <<https://leedsbeckett.on.worldcat.org/oclc/1289390963>>.

Urdan, T. C. (2022b) The Normal Distribution [Online]. In: *Statistics in Plain English*. Routledge, pp. Chapter 4. Available from: <<https://leedsbeckett.on.worldcat.org/oclc/1289390963>>.

Wichura, M. J. (1988) Algorithm AS 241: The Percentage Points of the Normal Distribution. *Journal of the Royal Statistical Society. Series C (Applied Statistics)* [Online], 37 (3), pp. 477–484. Available from: <<https://www.jstor.org/stable/2347330>> [Accessed 4 June 2025].